




















Tez — every measurement, one page. Zero-shot frozen Gemma 4 12B Q8 0 · one forward pass · RTX 5080 laptop (16 GB)

Laya checkpoints ran on the SAME rows and GPU; Jev figures are third-party published, never measured here.



Headline

typed-decisions zero-shot 0.704
(Laya base 0.36 · Jev 0.727 · Laya fine-tuned 0.766)

MASSIVE, 11 languages 0.885 macro
(Laya-multilingual 0.524 · Khmer 0.790 vs 0.000)

Banking77, 77 options 0.713 (Laya 0.395)

SemIf authored144 0.943 (SemIf 4B 0.813)

voice: per streamed word 30 ms compute · 1 harmful action / 198
order flip @ 20 options 0.07 (Laya 0.15-0.20 · Jev 0.13)

conformal 90% coverage act on 52% @ 0.853 accuracy

Where others win: Laya on tasks in its training mix
Jev on Banking77 (0.870).