
















Tez probe lab: reading decisions from the middle of a frozen model (typed-decisions, 2,000 test decisions, unless stated)



What carries over to any schema

- 1 read the decision at the layer where the unlabelled states are most compressed (intrinsic-dimension minimum: layers 24-28)
- 2 cut the model there: fewer layers, same probe
- 3 label the most typical rows first, then escalate, but keep a random audit slice
- 4 act only under a conformal error bound (5 %: acts on 40 % of decisions)
- 5 keep 64 numbers per state: the decision code

What did not work: label-free probes beyond their teacher, Dawid-Skene over weak readouts, uncertainty-only labelling, DoLa, adaptive depth, one shared letter probe on unseen workflows, a collapsing voice bound.